

# EndOfSemesterProject

June 5, 2024

## 1 End of Semester Project

We've done a lot this semester, you've learned a few things and now you're going to have fun (hopefully) making and thinking about some more advanced plotting.

First, we will introduce additional visualization techniques available to use through the plotly library. Understanding the details of these visualization methods is out of the scope of this course, but they're still pretty cool to look at!

Your end of semester project will be to use these techniques to analyze some data. (More instructions will be given below.)

```
[1]: import numpy as np
import plotly.express as px
import seaborn as sns
import pandas as pd
```

### 1.1 Animated Scatter Plots

Our data today comes from the [Gapminder Foundation](#) which explores data on poverty, inequality and health around the world.

```
[2]: world = px.data.gapminder()
world.head(10)
```

```
[2]:      country continent  year  lifeExp      pop  gdpPercap iso_alpha \
0  Afghanistan      Asia  1952   28.801   8425333   779.445314      AFG
1  Afghanistan      Asia  1957   30.332   9240934   820.853030      AFG
2  Afghanistan      Asia  1962   31.997  10267083   853.100710      AFG
3  Afghanistan      Asia  1967   34.020  11537966   836.197138      AFG
4  Afghanistan      Asia  1972   36.088  13079460   739.981106      AFG
5  Afghanistan      Asia  1977   38.438  14880372   786.113360      AFG
6  Afghanistan      Asia  1982   39.854  12881816   978.011439      AFG
7  Afghanistan      Asia  1987   40.822  13867957   852.395945      AFG
8  Afghanistan      Asia  1992   41.674  16317921   649.341395      AFG
9  Afghanistan      Asia  1997   41.763  22227415   635.341351      AFG

      iso_num
0           4
```


```

1      4
2      4
3      4
4      4
5      4
6      4
7      4
8      4
9      4

```

Use the code to generate an animated scatter plot of GDP per capita and life expectancy over time. You can play the animation or scroll to a specific year.

```
[3]: px.scatter(world,
    x = 'gdpPerCap',
    y = 'lifeExp',
    hover_name = 'country',
    color = 'continent',
    size = 'pop',
    size_max = 60,
    log_x = True,
    range_y = [30, 90],
    animation_frame = 'year',
    title = 'Life Expectancy, GDP Per Capita, and Population over Time'
)
```




## 1.2 Animated Histograms

We can do the same with `px.histogram`, using the optional argument `animation_frame`.

```
[4]: px.histogram(world,
    x = 'lifeExp',
    animation_frame = 'year',
    range_x = [20, 90],

```

```
range_y = [0, 50],
title = 'Distribution of Life Expectancy over Time')
```



### 1.3 Box Plots

Box plots, also called “box and whisker plots” show the rough distribution of multiple numerical variables. In particular, they show the 25th, 50th (median), and 75th percentiles (the box), as well as 1.5 times the Interquartile Range (the whiskers). This is helpful for identifying outliers.


```
[5]: worldLatest = world[world['year']== 2007]
worldLatest.head(10)
```

```
[5]:
```

	country	continent	year	lifeExp	pop	gdpPercap	iso_alpha \
11	Afghanistan	Asia	2007	43.828	31889923	974.580338	AFG
23	Albania	Europe	2007	76.423	3600523	5937.029526	ALB
35	Algeria	Africa	2007	72.301	33333216	6223.367465	DZA
47	Angola	Africa	2007	42.731	12420476	4797.231267	AGO
59	Argentina	Americas	2007	75.320	40301927	12779.379640	ARG
71	Australia	Oceania	2007	81.235	20434176	34435.367440	AUS
83	Austria	Europe	2007	79.829	8199783	36126.492700	AUT
95	Bahrain	Asia	2007	75.635	708573	29796.048340	BHR
107	Bangladesh	Asia	2007	64.062	150448339	1391.253792	BGD
119	Belgium	Europe	2007	79.441	10392226	33692.605080	BEL

```
iso_num
11      4
23      8
35     12
47     24
59     32
71     36
83     40
95     48
107    50
```

```
[6]: px.box(worldLatest,
           y = 'lifeExp',
           x = 'continent',
           color = 'continent',
           hover_name = 'country',
           title = 'Distribution of Life Expectancy in 2007 by Continent'
           )
```



## 1.4 Violin Plots

These are basically the same as box plots, but show the distribution as well

```
[7]: px.violin(worldLatest,
                y = 'lifeExp',
                x = 'continent',
                color = 'continent',
                hover_name = 'country',
                title = 'Distribution of Life Expectancy in 2007 by Continent'
                )
```



## 1.5 Pie Charts

Pie charts look cool visually, but they are often hard to analyze. Let's take a closer look at some of them.

```
[8]: worldLatestAmericas = worldLatest[worldLatest['continent']=='Americas']
worldLatestAmericas.head(10)
```

```
[8]:
```


	country	continent	year	lifeExp	pop	gdpPercap	\
59	Argentina	Americas	2007	75.320	40301927	12779.379640	
143	Bolivia	Americas	2007	65.554	9119152	3822.137084	
179	Brazil	Americas	2007	72.390	190010647	9065.800825	
251	Canada	Americas	2007	80.653	33390141	36319.235010	
287	Chile	Americas	2007	78.553	16284741	13171.638850	
311	Colombia	Americas	2007	72.889	44227550	7006.580419	
359	Costa Rica	Americas	2007	78.782	4133884	9645.061420	
395	Cuba	Americas	2007	78.273	11416987	8948.102923	
443	Dominican Republic	Americas	2007	72.235	9319622	6025.374752	
455	Ecuador	Americas	2007	74.994	13755680	6873.262326	

	iso_alpha	iso_num
59	ARG	32
143	BOL	68
179	BRA	76
251	CAN	124
287	CHL	152
311	COL	170
359	CRI	188
395	CUB	192
443	DOM	214
455	ECU	218

```
[9]: px.pie(worldLatestAmericas,
            values = 'pop',
            names = 'country',
            title = 'Population of the Americas'
        )
```

Population of the Americas



We can see the countries showing up as percentages, but there are so many countries, so it's a little bit harder.

Let's take a look at what happens if we group by continent instead.

```
[10]: # Grouping by continent and summing up populations
worldLatestByContinent= worldLatest.groupby('continent')['pop'].sum().
    ↪reset_index()
worldLatestByContinent.head(10)
```

```
[10]:   continent      pop
0    Africa  929539692
1  Americas  898871184
2     Asia  3811953827
3   Europe  586098529
4  Oceania   24549947
```

```
[11]: px.pie(worldLatestByContinent,
            values = 'pop',
            names = 'continent',
            title = 'World Population by Continent')
```

World Population by Continent



## 1.6 Animated Pie Charts

This is pretty cool, but do you think we can animate the pie chart? Of course we can!!

First - let's go back to our original `world` DataFrame and get the population for each continent by year. We will use `group` and `retain` continent and year.

```
[12]: worldByContinent = world.groupby(['continent', 'year'])['pop'].sum().
    ↪reset_index()
worldByContinent.head(10)
```

```
[12]: continent year      pop
0    Africa 1952 237640501
1    Africa 1957 264837738
2    Africa 1962 296516865
3    Africa 1967 335289489
4    Africa 1972 379879541
5    Africa 1977 433061021
6    Africa 1982 499348587
7    Africa 1987 574834110
8    Africa 1992 659081517
9    Africa 1997 743832984
```

Now we will do the same field we did before in order to animate by year.

However, because there doesn't seem to be the same animate frames we had to do something a little stranger. Create a pie chart for each year, and then animate it with a figure.

It's a little clunkier, but it seems to work.

```
[13]: import plotly.graph_objs as go

# Assuming worldByContinent already exists

# Create pie charts for each year
frames = []
for year in worldByContinent['year'].unique():
    data_year = worldByContinent[worldByContinent['year'] == year]
    pie = go.Pie(labels=data_year['continent'], values=data_year['pop'],
    ↪name=f'Year {year}')
    frame = go.Frame(data=[pie], name=f'Year {year}')
    frames.append(frame)

# Create figure with the first year's pie chart
fig = go.Figure(data=[go.Pie(labels=worldByContinent[worldByContinent['year']
    ↪== min(worldByContinent['year']))['continent'],
    ↪values=worldByContinent[worldByContinent['year']
    ↪== min(worldByContinent['year']))['pop'],
    ↪name=f'Year {min(worldByContinent["year"])}']),
    ↪frames=[frames[0]])

# Add frames to animation
fig.frames = frames

# Update layout
fig.update_layout(title=f'World Population by Continent - Year
    ↪{min(worldByContinent["year"])}',
    ↪updatemenus=[{
    ↪    'type': 'buttons',
```

```


        'buttons': [{
            'label': 'Play',
            'method': 'animate',
            'args': [None, {
                'frame': {
                    'duration': 1000, # Set the duration for
↪ each frame (milliseconds)
                    'redraw': True
                },
                'fromcurrent': True,
                'mode': 'immediate'
            }]
        }, {
            'label': 'Pause',
            'method': 'animate',
            'args': [[None], {
                'frame': {
                    'duration': 0,
                    'redraw': False
                },
                'mode': 'immediate'
            }]
        }]
    })

# Update each frame's title to include the corresponding year
for i, frame in enumerate(fig.frames):
    fig.frames[i].update(layout_title_text=f'World Population by Continent -
↪ Year {worldByContinent["year"].unique()[i]}')

fig.show()

```

World Population by Continent - Year 1952



Now, here is a moment that we ask ourselves, was it worth it?

The percentages per continent do not seem to change very much. Perhaps a line graph is a better



way to show this.

## 1.7 Line Graph

As mentioned before, pie charts are kind of hard to read sometimes. And looking at an animated pie chart like we did above was hard to read. Let's go back to one of the first types of plots we looked at here.

Line plots!!

```
[14]: worldByContinent['continent'].unique()
```

```
[14]: array(['Africa', 'Americas', 'Asia', 'Europe', 'Oceania'], dtype=object)
```


```
[15]: # Define the continents
continents = ['Africa', 'Americas', 'Asia', 'Europe', 'Oceania']

# Create a line graph for each continent
data = []
for continent in continents:
    continent_data = worldByContinent[worldByContinent['continent'] ==
↪continent]
    data.append(go.Scatter(x=continent_data['year'],
                           y=continent_data['pop'],
                           mode='lines',
                           name=continent))

# Create figure
fig = go.Figure(data=data)

# Update layout
fig.update_layout(title='Population of Each Continent Over Time',
                   xaxis_title='Year',
                   yaxis_title='Population')

fig.show()
```



This gets us some of the way there, but let's look into the percentage of each continent (like our pie chart from above).

Let's first add a column to give the percentage per year that each continent's population has.

```
[16]: total_population_by_year = worldByContinent.groupby('year')['pop'].sum()

worldByContinent['pop_pct'] = worldByContinent.apply(lambda row: row['pop'] /
    ↪total_population_by_year[row['year']], axis=1)

worldByContinent.head(10)
```

```
[16]:  continent  year      pop  pop_pct
0    Africa  1952  237640501  0.098731
1    Africa  1957  264837738  0.099398
2    Africa  1962  296516865  0.102255
3    Africa  1967  335289489  0.104209
4    Africa  1972  379879541  0.106201
5    Africa  1977  433061021  0.110192
6    Africa  1982  499348587  0.116414
7    Africa  1987  574834110  0.122527
8    Africa  1992  659081517  0.128961
9    Africa  1997  743832984  0.134870
```

And now let's plot the proportion of each continent by year.


```
[17]: # Define the continents
continents = ['Africa', 'Americas', 'Asia', 'Europe', 'Oceania']

# Create a line graph for each continent
data = []
for continent in continents:
    continent_data = worldByContinent[worldByContinent['continent'] ==
    ↪continent].copy()
    data.append(go.Scatter(x=continent_data['year'],
                           y=continent_data['pop_pct'],
                           mode='lines',
                           name=continent))

# Create figure
fig = go.Figure(data=data)

# Update layout
fig.update_layout(title='Percentage of Population by Continent Over Time',
                   xaxis_title='Year',
                   yaxis_title='Percentage of Population')
```

```
fig.show()
```



## 1.8 Timelines (Gantt Charts)

You may - or may not - have seen these before, but you can illustrate a project timeline with a Gantt chart.

Below this Gantt chart illustrates the timeline of the Suraj Rampure [Suraj Rampure](#), creator of Data 6 at UC Berkeley. (Which we have used for many of the ideas in this course!)

```
[18]: # Define the phases list
phases = [
    ['Newborn', '1998-11-26', '1999-11-26'],
    ['Toddler, Preschooler', '1999-11-26', '2005-09-03'],
    ['Elementary School Student', '2005-09-03', '2009-06-30'],
    ['Middle School Student', '2009-09-15', '2012-06-15'],
    ['High School Student', '2012-09-05', '2016-05-30'],
    ['Undergrad @ UC Berkeley', '2016-08-22', '2020-05-15'],
    ['Masters @ UC Berkeley', '2020-08-25', '2021-05-14'],
    ['Teaching Data 94', '2021-01-20', '2021-05-14']
]

# Define column labels
columns = ['Phase', 'Start', 'End']

# Create a DataFrame from the phases list
phases_df = pd.DataFrame(phases, columns=columns)

# Convert 'Start' and 'End' columns to datetime
phases_df['Start'] = pd.to_datetime(phases_df['Start'])
phases_df['End'] = pd.to_datetime(phases_df['End'])

phases_df.head(10)
```

```
[18]:
```

	Phase	Start	End
0	Newborn	1998-11-26	1999-11-26
1	Toddler, Preschooler	1999-11-26	2005-09-03
2	Elementary School Student	2005-09-03	2009-06-30
3	Middle School Student	2009-09-15	2012-06-15
4	High School Student	2012-09-05	2016-05-30
5	Undergrad @ UC Berkeley	2016-08-22	2020-05-15
6	Masters @ UC Berkeley	2020-08-25	2021-05-14
7	Teaching Data 94	2021-01-20	2021-05-14

```
[ ]: # Create a Gantt chart
fig = px.timeline(phases_df, x_start='Start', x_end='End', y='Phase')

# Update layout
fig.update_layout(title='Phases Timeline',
                  xaxis_title='Date',
                  yaxis_title='Phase')

fig.show()
```

Notice that the lifetime is increasing for the Gantt chart. You can also repeat things in a Gantt chart. For example here is - according to [Chat Gpt](#) - the Gantt chart for locations/scenes in *Star Wars IV: A New Hope*

```
[ ]: # Define the scenes and their corresponding start and end times in the original
      ↪ Star Wars movie
scenes = [
    ['Tatooine', '00:00:00', '00:10:00'],
    ['Death Star', '00:10:00', '00:15:00'],
    ['Tatooine', '00:15:00', '00:20:00'],
    ['Tatooine', '00:20:00', '00:25:00'],
    ['Death Star', '00:25:00', '00:30:00'],
    ['Yavin 4', '00:30:00', '00:35:00'],
    ['Death Star', '00:35:00', '00:40:00'],
    ['Tatooine', '00:40:00', '00:45:00'],
    ['Yavin 4', '00:45:00', '01:00:00'],
    ['Death Star', '01:00:00', '01:05:00'],
    ['Yavin 4', '01:05:00', '01:15:00'],
    ['Death Star', '01:15:00', '01:20:00'],
    ['Yavin 4', '01:20:00', '01:25:00']
]

# Convert to DataFrame
scenes_df = pd.DataFrame(scenes, columns=['Location', 'Start', 'End'])

# Convert 'Start' and 'End' columns to datetime
scenes_df['Start'] = pd.to_datetime(scenes_df['Start'], format='%H:%M:%S')
```

```

scenes_df['End'] = pd.to_datetime(scenes_df['End'], format='%H:%M:%S')

# Calculate duration of each scene
scenes_df['Duration'] = scenes_df['End'] - scenes_df['Start']

# Create a Gantt chart
fig = px.timeline(scenes_df, x_start='Start', x_end='End', y='Location',
                 title='Star Wars: A New Hope Scene Locations Timeline', hover_data=None)

# Update x-axis tick format to show only time component
fig.update_xaxes(tickformat='%H:%M:%S')

# Update layout
fig.update_layout(xaxis_title='Time', yaxis_title='Location')

fig.show()

```

## 1.9 Choropleth Maps

We have already seen Choropleth Maps, but they are super fun!

Let's go back to `worldLatest` which has the population information for 2007.

```
[20]: worldLatest.head(10)
```

```

[20]:
   country continent  year  lifeExp      pop      gdpPercap iso_alpha \
11  Afghanistan   Asia  2007   43.828  31889923    974.580338    AFG
23   Albania     Europe  2007   76.423   3600523   5937.029526    ALB
35   Algeria     Africa  2007   72.301  33333216   6223.367465    DZA
47   Angola     Africa  2007   42.731  12420476   4797.231267    AGO
59  Argentina  Americas  2007   75.320  40301927  12779.379640    ARG
71  Australia  Oceania   2007   81.235  20434176  34435.367440    AUS
83   Austria     Europe  2007   79.829   8199783  36126.492700    AUT
95   Bahrain     Asia   2007   75.635    708573  29796.048340    BHR
107  Bangladesh   Asia   2007   64.062  150448339   1391.253792    BGD
119   Belgium     Europe  2007   79.441  10392226  33692.605080    BEL

```


```

iso_num
11      4
23      8
35     12
47     24
59     32
71     36
83     40
95     48
107    50
119    56

```

```
[21]: px.choropleth(worldLatest,
                    locations = 'iso_alpha',
                    color = 'lifeExp',
                    hover_name = 'country',
                    title = 'Life Expectancy Per Country',
                    color_continuous_scale = px.colors.sequential.tempo
                    )
```

Life Expectancy Per Country




## 1.10 Animated Choropleth Maps

And of course, our course wouldn't be complete without animating these maps!

```
[22]: px.choropleth(world,
                    locations="iso_alpha",
                    color="lifeExp",
                    animation_frame="year",
                    color_continuous_scale = px.colors.diverging.RdYlGn,
                    title = "Life Expectancy Over Time",
                    range_color=(30,80))
```

Life Expectancy Over Time



### 1.11 3D Scatter Plots

It is also possible to plot points along three dimensions (i.e. with three coordinates).

Let's use a dataset that's part of seaborn, that has recorded data on a number of different penguins. They have recorded the: - Species - Island (they were observed on) - Length of Bill (in millimeters) - Depth of Bill (in millimeters) - Flipper Length (in millimeters) - Body Mass (in grams) - Sex (when possible to determine! Some missing values.)

```
[23]: penguins = sns.load_dataset('penguins')
      penguins.head(10)
```

```
[23]:  species      island  bill_length_mm  bill_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen         39.1           18.7             181.0
1  Adelie  Torgersen         39.5           17.4             186.0
2  Adelie  Torgersen         40.3           18.0             195.0
3  Adelie  Torgersen          NaN           NaN              NaN
4  Adelie  Torgersen         36.7           19.3             193.0
5  Adelie  Torgersen         39.3           20.6             190.0
6  Adelie  Torgersen         38.9           17.8             181.0
7  Adelie  Torgersen         39.2           19.6             195.0
8  Adelie  Torgersen         34.1           18.1             193.0
9  Adelie  Torgersen         42.0           20.2             190.0

      body_mass_g      sex
0         3750.0    Male
1         3800.0  Female
2         3250.0  Female
3            NaN     NaN
4         3450.0  Female
5         3650.0    Male
6         3625.0  Female
7         4675.0    Male
8         3475.0     NaN
9         4250.0     NaN
```

Let's plot each penguin in 3D space by their: - bill length - bill depth - flipper length

Let's use different colors to indicate their species.

Try dragging the, graph to move around the camera.


```
[24]: px.scatter_3d(penguins,
                  x = 'bill_length_mm',
                  y = 'bill_depth_mm',
                  z = 'flipper_length_mm',
                  color = 'species',
                  hover_name = 'island',
                  title = 'Flipper Length vs. Bill Depth vs. Bill Length')
```

species  
 • Adelie  
 • Chinstrap  
 • Gentoo

It is generally hard to work with plots in 3D. So often you need to consider is this really the best way to show the data?

Let's take a look at a 2D scatter plot and see if it's more or less informative.

```
[25]: px.scatter(penguins,
                x = 'bill_length_mm',
                y = 'bill_depth_mm',
                color = 'species',
                hover_name = 'island',
                title = 'Bill Depth vs. Bill Length')
```



What is something you notice about the different species?

If you had a new penguin with bill length of 45 mm and bill depth of 14 mm, could you make a guess as to what that species would be?

## 2 Back to the End of Semester Project!!

Okay - so we've had a lot of fun! Now it's time for you to get to work on your own!

**Instructions:** Your job for the End of Semester Project is to create **4 different plots using Python** (each with a different plot type) and provide a **short description (3 - 5 sentences)**



describing what this plot illustrates to you.

- **What to Submit:** You will need to turn in both a PDF; your ipynb and any data sets you use (you can provide links if they are online or upload the file itself.)
- **What Data Can I Use?** You can choose whatever data set you want for the plots. You can use the same data set for all 4 or different data sets for each. You can use any data set we have used in class, any of the built in python data sets (we will give some examples below) or (as in the example of Suraj Rampure) the data set may come from a personal place. (See below for some examples!)
- **What Kinds of Plots Can I Use?** You can use any type of plot that you want! You just need to be able to make it in Python and provide your code. Make sure your plot is readable.
- **Grading:** This will be worth 40 points in total, 10 points per plot split equally between the plot itself and your written description of it.

## 2.1 Plot Types:

Feel free to use any of the plot types we have used in class. Here's the ones we have used today.

- Scatter Plot (2 and 3D)
- Line Plot
- Pie Chart
- Histograms
- Box Plots
- Violin Plots
- Choropleth

Your plot can be animated (which would be super cool) but definitely doesn't need to be.

You can even use a plot type we haven't discussed in class such as a [Ternary Plots](#)

## 2.2 Data Sets

You can use whatever data sets you would like! We've used a lot of them in class (Zillow, UC Admission Data, Spotify, Fast Food).

### 2.2.1 Data Sets from Plotly

But Python itself has a number of really great built in data sets! Here are a few examples from `plotly.express` which we have imported as `px`

1 **Iris Dataset:** > A classic dataset in machine learning and statistics, containing measurements of iris flowers. > > `iris_df = px.data.iris()`

2 **Wind Dataset** > Contains wind speed and direction measurements collected from a meteorological station. > > `wind_df = px.data.wind()`

3 **Carshare Dataset** > Contains data on car sharing usage. > > `carshare_df = px.data.carshare()`

4 **Election Dataset** > Contains data on the 2008 US presidential election results. > > `election_df = px.data.election()`

```
[26]: #Iris Data Set
iris_df = px.data.iris()

#Wind Dataset:
wind_df = px.data.wind()

#Carshare Dataset:
carshare_df = px.data.carshare()

#Election Dataset:
election_df = px.data.election()
```

## 2.2.2 Data Sets from Seaborn

We have also seen other built in datasets from Seaborn. Here are some of them that you might also find helpful.

To load seaborn, you'll need to run `import seaborn as sns`

1. **Titanic Dataset:** > Contains data on passengers aboard the Titanic, including information on their survival status, age, sex, class, fare, and more. > > `titanic_df = sns.load_dataset('titanic')`
2. **Functional Magnetic Resonance Image (fMRI) FMRI Dataset:** > Contains functional magnetic resonance imaging (fMRI) data, with time series measurements taken at different points in the brain. > > `fmri_df = sns.load_dataset('fmri')`
3. **Diamonds Dataset:** > Contains data on diamonds, including their carat weight, cut, color, clarity, and price. > > `diamonds_df = sns.load_dataset('diamonds')`
4. **Exercise Dataset:** > Contains data from an exercise physiology study, including measurements such as subject IDs, age, weight, and the amount of exercise they performed. > > `exercise_df = sns.load_dataset('exercise')`

```
[27]: titanic_df = sns.load_dataset('titanic')

fmri_df = sns.load_dataset('fmri')

diamonds_df = sns.load_dataset('diamonds')

exercise_df = sns.load_dataset('exercise')
```

```
[ ]:
```